

tldraw-sync-gcp

Production-Ready Real-Time Collaboration
on Google Cloud Platform

Architecture & Performance Overview

tldraw + Google Cloud

Executive Summary

A horizontally scalable sync backend for tldraw, enabling real-time collaboration with **7,000+ concurrent connections** and **99.99% reliability**.

7K+

Concurrent
Users

99.99%

Success Rate

235ms

Connection
Latency

100+

Concurrent
Rooms

Key Achievements



Zero Split-Brain

Redis distributed locking ensures each room has exactly one owner at any time.



Seamless Scaling

Two-phase handover protocol enables safe room migration during pod scaling.



Zero Data Loss


GCS persistence with automatic snapshots ensures durability across restarts.



Low Latency

WebSocket-based sync with server-side keep-alive for reliable connections.

Architecture Overview



Key Design Decision: Consistent Hashing

NGINX routes requests by URI hash, ensuring all users in the same room connect to the same pod. When pods scale, the two-phase handover protocol safely migrates rooms.

Technology Stack

Backend Services

Component	Technology
Runtime	Node.js 20
HTTP Server	Express 5
WebSocket	ws library
Sync Protocol	@tldraw/sync-core
Metrics	Prometheus

GCP Services

Service	Purpose
GKE	Container orchestration
Cloud Storage	Room snapshots & assets
Memorystore	Redis for locking
Cloud Load Balancing	Traffic distribution
Workload Identity	Secure GCS access


Key Configuration Parameters

Parameter	Value	Purpose
LOCK_TIMEOUT_SEC	10s	Redis lock TTL (auto-expires if pod crashes)
HEARTBEAT_INTERVAL_MS	5s	Lock renewal frequency
PING_INTERVAL_MS	25s	WebSocket keep-alive (under GCP's 30s timeout)
HANDOVER_TIMEOUT_MS	5s	Max wait for lock release during migration
THROTTLE_SAVE_MS	10s	GCS snapshot frequency


Stateful Room Ownership

Each room is owned by exactly **one pod** at any time, preventing split-brain scenarios and ensuring data consistency across all connected users.

How It Works



Redis Lock Design



Failure Handling

Scenario	Behavior	Recovery Time
Pod crashes	Lock expires, another pod acquires	~10 seconds
Network partition	Lock expires, GCS has last saved state	~10 seconds
Graceful shutdown	Lock deleted immediately, snapshot saved	Instant

Two-Phase Coordinated Handover

When pods scale, the hash ring changes and rooms need to migrate. The two-phase protocol ensures users are **only disconnected after the new owner is ready**.

Why Two-Phase?

Problem: Single-phase handover closes user connections before the new pod is ready. If the new pod is slow or crashes, users get stuck in retry loops.

Solution: Wait for "ready" signal before disconnecting users.

Phase 1: Lock Transfer

1. Pod C fails to acquire lock (Pod A owns it)
2. Pod C subscribes to `handover-lock-released:{roomId}`
3. Pod C publishes request to `room-handover`
4. Pod A saves state to GCS
5. Pod A releases lock and signals

Phase 2: Ready Confirmation

1. Pod C acquires lock
2. Pod C loads state from GCS
3. Pod C signals `handover-ready:{roomId}`
4. Pod A receives signal, NOW closes sockets (1013)
5. Users reconnect to Pod C (ready!)

Redis Pub/Sub Channels


Channel	Direction	Purpose
<code>room-handover</code>	Pod C → Pod A	Request room release
<code>handover-lock-released:{roomId}</code>	Pod A → Pod C	Lock is free, acquire now

handover-ready:{roomId}

Pod C → Pod A

Ready to serve, close your sockets

Handover Sequence



Key Insight

Users stay connected to Pod A until step 9. They only disconnect when Pod C is fully ready to serve them. This eliminates the "retry loop" problem of single-phase handover.

Performance & Load Testing

Stress tested with k6 from within GCP (Cloud Shell) to eliminate network variables.

Test Infrastructure

NGINX Ingress

5 replicas

App Pods

10 replicas

GKE Nodes

3× e2-medium

Capacity Results

Concurrent Users	Configuration	Success Rate	Status
5,000	50 rooms × 100 users	100%	✓ Perfect
7,000	100 rooms × 70 users	99.99%	✓ Optimal
10,000	100 rooms × 100 users	30%	⚠ Exceeded capacity

Connection Latency

Scenario	Latency	Notes
Normal load	~235ms	Optimal user experience
Heavy load (7,000 VUs)	~10-16s	Connection queueing under stress
Peak load	~20s p(95)	At capacity limits

Scaling Recommendations

Current Capacity

~7,000 concurrent WebSocket connections with 99.99% reliability

To Scale Beyond 7,000 Connections

- 1. Increase NGINX Ingress Replicas**
Scale from 5 to 10+ replicas. Each NGINX pod handles ~1,400 connections.
- 2. Use Larger Node Types**
Upgrade from e2-medium to e2-standard-4 or larger for more CPU/memory per node.
- 3. Tune NGINX Worker Connections**
Increase `worker-connections` in NGINX ConfigMap (default: 1024 → 65536).
- 4. Add More Nodes**
Scale node pool from 3 to 5+ nodes for better pod distribution.

Estimated Capacity by Configuration

NGINX Replicas	App Pods	Nodes	Estimated Capacity
5	10	3× e2-medium	~7,000 connections
10	20	5× e2-standard-4	~15,000 connections
20	40	10× e2-standard-4	~30,000 connections

Monitoring & Observability

Prometheus Metrics

Metric	Type	Description
tldraw_active_rooms	Gauge	Rooms currently in memory
tldraw_active_connections	Gauge	Active WebSocket connections
tldraw_handover_requests_total	Counter	Handover requests initiated
tldraw_handover_success_total	Counter	Successful handovers completed
tldraw_handover_timeouts_total	Counter	Handover timeouts (indicates issues)
tldraw_handover_duration_seconds	Histogram	Time from request to lock acquisition

Recommended Alerts

High Handover Timeout Rate

```
rate(tldraw_handover_timeouts_total[5m])
/
rate(tldraw_handover_requests_total[5m])
> 0.1
```

May indicate network issues or pod health problems

Slow Handovers

```
histogram_quantile(0.95,
tldraw_handover_duration_seconds_bucket)
> 3
```

Handovers taking longer than expected

Health Endpoints

Endpoint	Purpose	Used By
GET /api/health	Liveness check	Kubernetes probes

GET /metrics

Prometheus metrics

Monitoring stack

Code Reference: Connection Flow

1. WebSocket Upgrade & Keep-Alive

```
// src/index.ts - WebSocket upgrade handling
server.on("upgrade", (request, socket, head) => {
  const { pathname, searchParams } = new URL(request.url, `http://${request.headers.host}`);
  const roomMatch = pathname.match(/\/api\/connect\/(.*)/);
  const roomId = roomMatch?.[1];
  const sessionId = searchParams.get("sessionId");

  if (roomMatch && roomId && sessionId) {
    wss.handleUpgrade(request, socket, head, (ws) => {
      wss.emit("connection", ws, request, roomId, sessionId);
    });
  }
});

// Keep-alive ping to prevent GCP LB idle timeout (30s)
const PING_INTERVAL_MS = 25000;

wss.on("connection", (ws, request, roomId, sessionId) => {
  const pingInterval = setInterval(() => {
    if (ws.readyState === WebSocket.OPEN) ws.ping();
  }, PING_INTERVAL_MS);

  ws.on("close", () => clearInterval(pingInterval));
  roomManager.getOrCreateRoom(roomId, ws, sessionId);
});
```

2. Redis Lock Acquisition

```
// src/roomManager.ts - Distributed lock with SETNX
const LOCK_TIMEOUT_SEC = 10;
const lockKey = `lock:room:${roomId}`;

// Try to acquire lock (atomic operation)
const lockAcquired = await redisClient.set(lockKey, POD_NAME, {
  EX: LOCK_TIMEOUT_SEC, // 10 second TTL
  NX: true,             // Only if Not eXists
});

if (lockAcquired) {
  // We own the room - start heartbeat to maintain lock
  const lockHeartbeat = setInterval(() => {
    redisClient.set(lockKey, POD_NAME, {
      EX: LOCK_TIMEOUT_SEC,
    });
  }, LOCK_TIMEOUT_SEC);
}
```

```
        XX: true, // Only if exists (we still own it)
    });
}, HEARTBEAT_INTERVAL_MS); // Every 5 seconds
}
```


Code Reference: Two-Phase Handover

3. Phase 1: Request & Release Lock

```
// Pod C (new target) - Request handover
const subscription = await this.subscribeToLockReleased(roomId);

await pubClient.publish("room-handover", JSON.stringify({
  roomId,
  targetPodId: POD_NAME,
}));

const lockReleased = await subscription.wait(); // Wait up to 5s

// Pod A (current owner) - Handle handover request
await subClient.subscribe("room-handover", async (message) => {
  const request = JSON.parse(message);
  if (this.activeRooms.has(request.roomId)) {
    await this.releaseRoom(request.roomId);
  }
});
```

4. Phase 2: Ready Signal & Migration

```
// Pod A - Release room and wait for ready signal
async releaseRoom(roomId) {
  // 1. Save state to GCS
  const snapshot = room.getCurrentSnapshot();
  await persistRoomSnapshot(roomId, snapshot);

  // 2. Release Redis lock
  await redisClient.del(`lock:room:${roomId}`);

  // 3. Signal lock released, then wait for ready
  await pubClient.publish(`handover-lock-released:${roomId}`, {...});
  const isReady = await this.waitForReadySignal(`handover-ready:${roomId}`);

  // 4. NOW close user connections (they'll reconnect to Pod C)
  for (const ws of sockets) {
    ws.close(1013, "Room migrated to another server, please reconnect");
  }
}

// Pod C - Acquire lock and signal ready
const { acquired, shouldSignalReady } = await this.acquireLockWithHandover(roomId);
if (acquired && shouldSignalReady) {
```

```
await pubClient.publish(`handover-ready:${roomId}`, { newOwner: POD_NAME });  
}
```

Code Reference: Persistence & Shutdown

5. GCS Persistence (Throttled)

```
// src/gcsStorage.ts - Room snapshot persistence
export async function persistRoomSnapshot(roomId: string, snapshot: RoomSnapshot) {
  const file = bucket.file(`rooms/${roomId}`);

  await retryOperation(async () => {
    await file.save(JSON.stringify(snapshot), {
      contentType: "application/json",
      resumable: false,
    });
  }, 3, 1000); // 3 retries, exponential backoff
}

// src/roomManager.ts - Throttled saves on data change
const saveToGCSThrottled = throttle(() => {
  const snapshot = room.getCurrentSnapshot();
  if (snapshot) persistRoomSnapshot(roomId, snapshot);
}, THROTTLE_SAVE_MS); // Max once per 10 seconds

room = new TLWebSocketRoom({
  schema,
  initialSnapshot: await fetchRoomSnapshot(roomId),
  onDataChange: () => saveToGCSThrottled(),
});
```

6. Graceful Shutdown

```
// src/roomManager.ts - Clean shutdown on SIGTERM
public async shutdown() {
  console.log(`Shutting down, saving ${this.activeRooms.size} rooms...`);

  for (const [roomId, room] of this.activeRooms) {
    // Stop heartbeat
    clearInterval(this.roomHeartbeats.get(roomId));

    // Save final snapshot
    const snapshot = room.getCurrentSnapshot();
    if (snapshot) await persistRoomSnapshot(roomId, snapshot);

    // Release lock immediately (so other pods can acquire)
    await redisClient.del(`lock:room:${roomId}`);
  }

  // Close all Redis connections
  await Promise.all([
```

```
        redisClient.quit(), subClient.quit(),  
        pubClient.quit(), handoverSubClient.quit()  
    ]);  
}
```

Summary

What We Built



Production-Ready Architecture

Horizontally scalable sync backend with stateful room ownership and distributed locking.



Safe Scaling

Two-phase coordinated handover ensures zero user disruption during pod scaling events.



Proven Performance

Stress tested to 7,000+ concurrent connections with 99.99% success rate.



GCP Native

Leverages GKE, Cloud Storage, Memorystore, and Workload Identity for secure, managed infrastructure.

Key Differentiators

- **Zero split-brain:** Redis locking guarantees single room ownership
- **Zero data loss:** GCS persistence with throttled snapshots
- **Seamless migration:** Two-phase handover with ready confirmation
- **Connection reliability:** Server-side keep-alive prevents idle timeouts

Documentation

Document	Description
<code>docs/architecture.md</code>	Comprehensive technical architecture
<code>docs/coordinated-handover.md</code>	Two-phase handover protocol details
<code>stress-test/README.md</code>	Load testing guide and results

Questions?

Repository: `github.com/tldraw/tldraw-sync-gcp`